

PXRD-Diff: diagnosing the lattice-recovery bottleneck in diffusion-based powder-diffraction structure inversion

Frank Cai

Purdue University, West Lafayette, IN, USA

frankyc11223@gmail.com

Manuscript type: Research Paper (Journal of Applied Crystallography).

Synopsis: A deliberately small conditional diffusion model for powder-diffraction structure inversion is used as a diagnostic instrument. A controlled ablation shows that restoring absolute peak positions to the lattice head, by position-aware pooling or by explicit d-spacings, does not reduce its ~ 1.3 Å cell error; the evidence points to the regression formulation of lattice recovery, not to the pooling stage. A classical autoindexer, supplied with the true crystal system, serves as a control; three external generative models are re-scored on one shared evaluation harness; and a code audit prompted by review corrects three implementation errors. Retraining with the corrected loss lifts the learned head from 0 to 1.9 %, level with the control, while the true-lattice oracle stays at 5.5 %, so the gap to the oracle is the effect that replicates.

Abstract

Inverting a one-dimensional powder X-ray diffraction (PXRD) pattern to a three-dimensional crystal structure remains an open problem. Recent generative approaches report results under mutually incompatible protocols, and their failure modes are seldom catalogued. This paper is a diagnostic study. PXRD-Diff, a 3.7 M-parameter conditional diffusion model trained on the CDVAE MP-20 split, is kept deliberately small so that its failure modes are visible; it is not proposed as a competitive inversion tool. DiffractGPT, PXRDnet and deCIFer are re-scored through the same pymatgen StructureMatcher harness and reach 18.9 %, 30.0 % and 73.8 % match respectively; PXRD-Diff reaches 1.9 %. The central finding is mechanistic. The learned PXRD encoder reaches a 0.007 normalised auxiliary lattice-regression loss in isolation, yet substituting that head’s own lattice into the sampler recovers only 1.2 % of structures, because its ~ 1.1 Å length error shifts every Bragg peak by ~ 15 %. Stronger, constrained supervision of the head does not reduce that error. A controlled ablation on the frozen encoder shows that restoring absolute peak positions to the head, by position-aware attention pooling or by explicit d-spacing features, leaves the error unchanged (1.28–1.31 Å across three arms, against 1.32 Å for the checkpoint’s own head), so the earlier reading that global pooling discards the d-spacings is withdrawn. The evidence now points to the regression formulation itself: a head handed the same d-spacings the autoindexer reads does not perform the peak-to-index assignment that turns them into a cell. A perturbation study locates

a sharp sensitivity knee at ~ 0.5 Å cell error. A true-lattice oracle lifts match to 5.5 % at $n = 1000$ (three seeds; 95 % Wilson confidence interval [4.7, 6.4]), against 1.9 % [1.5, 2.5] for the learned head (paired McNemar $p < 10^{-1}$). As a classical control, a Q-space autoindexer that is supplied with the true crystal system replaces the learned lattice head and reaches 1.6 % [1.2, 2.1] (1.4 % without the crystal system); its cells are far closer to the truth (pattern Pearson 0.41 against 0.18), yet at three seeds it ties the learned head on match rate (42 fixed, 52 broken, $p = 0.35$), fixing hexagonal and tetragonal structures and losing cubic ones. A code audit prompted by review found and corrected three implementation errors, including an atomic form factor evaluated in the wrong convention and a Debye loss evaluated on the wrong tensor. The corrected simulator agrees with pymatgen at a median Pearson correlation of 0.999 over 1000 structures. Retraining the production checkpoint with the corrections raised the learned head from 0 to 1.9 % and removed the indexer’s aggregate lead reported in the submitted version; the oracle gap is the result that replicates. The contributions are the lattice-recovery bottleneck diagnosis, a four-item failure catalogue, the shared evaluation harness and the audited, released code.

Keywords: powder diffraction; autoindexing; crystal structure determination; conditional diffusion models; machine learning; differentiable simulation; reproducibility

1. Introduction

PXRD is the most common structural-characterization measurement in solid-state chemistry. The forward problem is solved by 1920s structure-factor physics and a one-line `pymatgen` call. The inverse is brutal: symmetry-equivalent structures produce identical Bragg reflections, intensities are orientation-averaged, and the 3D $\rightarrow$ 1D reduction destroys atom labelling. The traditional indexing $\rightarrow$ Pawley/Le Bail $\rightarrow$ Rietveld pipeline needs a trained crystallographer and a starting model good enough to converge; for the long tail of new materials, structures stay unsolved.

Three recent generative approaches (DiffractGPT: Choudhary, 2025; Crystalyze: Riesel *et al.*, 2024; PXRDnet: Guo *et al.*, 2025) argue neural networks can shortcut this pipeline. They report match rates in the tens of percent, but under mutually incompatible evaluation protocols, and failure modes are seldom catalogued. A practitioner cannot easily form a calibrated expectation of *what is hard* about PXRD inversion.

We build a small reproducible conditional diffusion model (**PXRD-Diff**) and use it as a diagnostic instrument. Its size is a design choice: a 3.7 M-parameter model exposes a lattice-recovery bottleneck that capacity would hide. It is not proposed as an alternative to the larger systems it is compared with. Contributions:

1. **A small, fully released baseline.** 3.7 M parameters, trains on a single RTX 5090 in ~ 1.5 h on CDVAE MP-20. A PyTorch port of the `pymatgen` structure-factor calculation, differentiable with respect to coordinates and lattice, supplies an auxiliary pattern-matching loss. The port introduces no new physics (§3.3).
2. **Lattice-recovery bottleneck diagnosis (the primary contribution).** A perturbation study (§5.5) locates a sharp ~ 0.5 Å sensitivity knee. A true-lattice oracle lifts match to **5.5 %** at $n = 1000$ (three seeds; Wilson 95 % CI [4.7, 6.4]) from 1.9 % [1.5, 2.5] for the learned head (paired McNemar $p < 10^{-1}$). The encoder reaches a 0.007 *normalised* aux lattice loss, yet its head’s lattice recovers only 1.2 % end to end, a ~ 1.1 Å error well above the knee (§5.6). A controlled pooling ablation (§5.6) finds that neither position-aware pooling nor explicit

peak-position features reduce that error, so the evidence points to the regression formulation rather than to the pooled encoding. The 5.5 % oracle ceiling shows coordinates remain a co-bottleneck.

3. **A classical control.** A Q-space autoindexer, supplied with the true crystal system, replaces the learned lattice head at sampling time and reaches 1.6 % [1.2, 2.1]. With the submitted checkpoint the learned head recovered nothing and the indexer’s lead was significant (27 fixed, 0 broken); with the checkpoint retrained under the corrected loss the two tie at three seeds (42 fixed, 52 broken, $p = 0.35$), the indexer winning the hexagonal and tetragonal structures below the 0.5 Å knee and the learned head the cubic ones (§5.2). This is a sanity check on the diagnosis, not a method: a fifty-year-old algorithm given extra prior information is expected to match or beat the learned head. Without the crystal system the indexer indexes 43.4 % of cells against 48.8 % and recovers 1.4 % [1.0, 1.9] downstream (§3.5).
4. **Four-item failure catalogue.** Top-K Debye rerank, Debye-gradient DDIM guidance, Wyckoff-letter embeddings and a distance-matrix aux loss all fail at $n \geq 200$; the latter two combined (v14) underperform the bare ε baseline. The gradient-guidance failure agrees with independent work showing the PXRD loss landscape is too rough for gradient descent (Segal *et al.*, 2025). The Wyckoff result is read narrowly: the embedding carried a label, not a constraint (§5.4).
5. **Reproduced baselines and a shared harness.** DiffractGPT (Mistral-7B + LoRA), PXRDnet (CDVAE + XRD encoder) and deCIFer are re-scored on our harness: **18.9 %** match (DGpt, $n = 1000$), **30.0 %** (PXRDnet, $n = 20$) and **73.8 %** (deCIFer, $n = 298$). These numbers set the scale of the diagnostic model; they are not a claim of competitiveness. A candidate divergence between match rate and space-group recovery across two baselines is raised in §5.3 and §6 as a hypothesis that the $n = 20$ PXRDnet sample cannot establish.
6. **A code audit.** Review of the released code identified three implementation errors: the atomic form factor was evaluated in the wrong convention, the Debye loss in x_0 -residual mode was evaluated on the wrong tensor, and patterns the indexer could not index were silently given the true lattice. §5.7 reports each error, its effect on every published number, and what was re-run.

2. Related work

Diffusion for crystals. CDVAE (Xie *et al.*, 2022) introduced the MP-20 protocol; DiffCSP (Jiao *et al.*, 2023) used joint coord+lattice diffusion; MatterGen (Zeni *et al.*, 2025) scaled to millions of structures. PXRD-Diff borrows DiffCSP’s joint diffusion, restricted to the conditional setting where a 1D pattern conditions generation.

Neural PXRD inversion. DiffractGPT (Choudhary, 2025) is a fine-tuned Mistral-7B that decodes peak list + formula $\rightarrow$ CIF tokens. Crystalyze (Riesel *et al.*, 2024) uses CDVAE conditioned on an XRD transformer encoder with post-hoc symmetry filtering. PXRDnet (Guo *et al.*, 2025) is a CDVAE variant with *iterative latent-space gradient guidance* (~ 3.5 M decoder ops/material), published in *Nature Materials* in 2025. deCIFer (Johansen *et al.*, 2025) is a concurrent autoregressive-language-model approach that, like DiffractGPT, emits CIF tokens. All four are larger than PXRD-Diff and report results under bespoke evaluation pipelines; §5.3 re-scores DiffractGPT, PXRDnet, and deCIFer on a single shared **StructureMatcher** harness. Crystalyze’s checkpoint download link is “not yet active” (verified 2026-06-01).

Where the unit cell comes from. Recent approaches differ in whether the unit cell is generated jointly with the coordinates or supplied as a separate input. PXRDGen (Li *et al.*, 2025) generates lattice and coordinates jointly by default, but can accept an independently determined cell, from a cell network or from conventional indexing, as an additional condition. XtalNet (Lai *et al.*, 2025) is an end-to-end equivariant approach, reported on metal–organic-framework benchmarks. Parackal *et al.* (2024) enumerate symmetry-confined Wyckoff arrangements and rank them with a learned energy model before fitting the pattern. XRDSol (Yu *et al.*, 2026) denoises coordinates with an equivariant diffusion model conditioned on the pattern, and takes the stoichiometry and the unit cell as inputs. That last design choice is the one this paper examines from the other side: XRDSol assumes an indexed cell, and the cell is exactly what our learned head fails to recover. Chitturi *et al.* (2021) regress lattice parameters from simulated patterns with a convolutional network and analyse how impurities, noise and broadening degrade the prediction, a diagnostic aim close to ours. None is reproduced here; they place PXRD-Diff, a diagnostic instrument rather than a competing solver, within the range of current designs.

Classical structure solution from powder data. Recovering a unit cell from peak positions is a solved sub-problem with a fifty-year toolchain: ITO, TREOR (Werner *et al.*, 1985) and DICVOL (Boultif & Louër, 2004), ranking solutions with the de Wolff (1968) figure of merit M20, index high-symmetry powder patterns in CPU milliseconds, with the difficulty concentrated in low-symmetry (monoclinic and triclinic) cells. Indexing is only the first stage of classical structure solution. Once the cell and space group are fixed, direct-space search, simulated annealing and charge flipping complete the structure, as implemented in programs such as FOX (Favre-Nicolin & Černý, 2002) and EXPO (Cuocci *et al.*, 2022). §3.5 reuses only the indexing stage, as a control for the learned lattice head. The indexer is given the true crystal system, which classical practice must itself infer; §3.5 also reports the run without it. §5.2 and §5.5 show the resulting lift tracks the symmetry-dependent accuracy classical indexers have always had.

Differentiable physics. Scattering-pattern losses date to Rietveld. The module used here is a PyTorch port of the `pymatgen` structure-factor computation, differentiable with respect to coordinates and lattice so that it can serve as a diffusion-training loss term; it adds no physics beyond that computation (§3.3). Concurrent generative-inversion work conditions on differentiable spectral targets for amorphous and nanostructured systems (Guo & Schwalbe-Koda, 2026). Relevant to our negative results, Segal *et al.* (2026) show the powder-XRD similarity loss landscape is “too rough for gradient descent”, independent evidence for why our sampling-time Debye-gradient guidance fails (§5.4).

Equivariance. Not strictly E(3)-equivariant: we use periodic-distance RBFs (SchNet-style; Schütt *et al.*, 2017). A frozen MACE (Batatia *et al.*, 2022) encoder showed no measurable benefit at MP-20 scale in pilots and was deferred.

3. Method

3.1 Problem formulation

Let $\mathcal{C} = (\mathbf{F}, \mathbf{Z}, \mathbf{L})$ denote a crystal structure with fractional coordinates $\mathbf{F} \in [0, 1)^{N \times 3}$, atomic numbers $\mathbf{Z} \in \{1, \dots, 100\}^N$, and a 3×3 lattice matrix $\mathbf{L}$ parametrised by $(a, b, c, \alpha, \beta, \gamma) \in \mathbb{R}^6$. Let $\mathbf{p}(\mathcal{C}) \in \mathbb{R}^{4251}$ be the simulated Cu K α PXRD pattern on a fixed 2θ grid from 5° to 90° at 0.02° resolution, normalised to maximum intensity 1.

We are given $\mathbf{p}$ and $\mathbf{Z}$ at test time, and we wish to sample $(\mathbf{F}, \mathbf{L}) \sim p(\cdot | \mathbf{p}, \mathbf{Z})$. Note that the composition $\mathbf{Z}$ being known is a meaningful simplification (in practice one usually knows the chemistry from synthesis), but it is consistent with prior work on the same task.

3.2 Architecture overview

PXRD-Diff has three trained components: a PXRD encoder, a denoiser, and a small auxiliary lattice head used only during training.

PXRD encoder. A 1D ResNet of four blocks (channels $64 \rightarrow 128 \rightarrow 256 \rightarrow 256$, stride-2 downsampling, GroupNorm, SiLU) ingesting the standardised pattern. We expose two outputs: a global pooled vector $\mathbf{g} \in \mathbb{R}^{256}$ and a multi-resolution feature map $\mathbf{F}_{\text{pxrd}} \in \mathbb{R}^{L \times 256}$ obtained by 1×1 -projecting each block output to d_{model} and concatenating along the spatial axis. *The multi-resolution map is critical.* In an early run (referred to as `gpu_v4` in our logs) we conditioned the denoiser on $\mathbf{g}$ only via additive broadcast, and the coordinate loss never moved off the random baseline of 3.0; the auxiliary head told us the encoder was learning useful features (its loss dropped from 0.99 to 0.007), but the `AdaptiveAvgPool1d(1)` collapse destroyed all spectral structure before it could reach the denoiser.

Denoiser. A periodic-distance message-passing network with $L = 3$ layers and $d_{\text{model}} = 256$. Atom embeddings $\mathbf{h}_i^{(0)} = \text{Emb}(\mathbf{Z}_i) + W_{\text{coord}} \mathbf{F}_i^{(t)}$ are updated by alternating (i) message passing over RBF-encoded periodic Cartesian distances under minimum-image convention, with timestep FiLM conditioning, and (ii) cross-attention from atoms (queries) to multi-resolution PXRD features (keys/values). Two output heads predict per-atom coordinate noise $\epsilon_F \in \mathbb{R}^{N \times 3}$ and lattice noise $\epsilon_L \in \mathbb{R}^6$. The lattice head is a small MLP applied to $[\bar{\mathbf{h}}; \mathbf{g}; W_{\text{lat}} \ell^{(t)}; \mathbf{t}_{\text{cond}}]$: that is, the pooled atom features, the global PXRD embedding, **a projection of the noisy lattice itself**, and the timestep encoding (see §3.4).

Auxiliary head. A single-hidden-layer MLP that predicts $(a, b, c, \alpha, \beta, \gamma)$ from $\mathbf{g}$ alone, trained with an MSE loss against the (normalised) ground-truth lattice. Its purpose is purely diagnostic, to expose whether the encoder is learning any lattice-relevant representation, and its loss does not flow back into the denoiser.

3.3 Differentiable Bragg structure-factor loss

This module is an engineering component, not a physical contribution. It is a PyTorch port of the structure-factor calculation in `pymatgen.analysis.diffraction.xrd.XRDCalculator`, written so that gradients flow to fractional coordinates and lattice. We implement

$$F(hkl) = \sum_j f_j(s) \exp(-B_{\text{iso}} s^2) \exp[2\pi i (hx_j + ky_j + lz_j)],$$

with the atomic form factor in the convention `pymatgen` uses for its tabulated coefficients,

$$f_j(s) = Z_j - 41.78214 s^2 \sum_{k=1}^4 a_k \exp(-b_k s^2),$$

where (a_k, b_k) are read from `ATOMIC_SCATTERING_PARAMS`, Z_j is the atomic number, $s = \sin \theta / \lambda$, and $B_{\text{iso}} = 0.5 \text{ \AA}^2$ is a uniform isotropic temperature factor. The intensity at each reflection is $|F(hkl)|^2$ multiplied by the standard Lorentz-polarisation correction. We enumerate all (h, k, l) with

$|h|, |k|, |l| \leq 5$ (1330 reflections after excluding the origin), place each on the 2θ grid via Bragg’s law, and broadcast each reflection as a Gaussian peak with FWHM 0.1° to obtain a continuous, differentiable pattern on a 256-bin coarse grid. The module treats ideal powders only: no preferred orientation, peak asymmetry, instrument profile or background is modelled. It is a single `nn.Module`, vectorised over the batch dimension.

Correction made during review. The submitted version of this paper evaluated the form factor as the bare sum $\sum_k a_k \exp(-b_k s^2)$ with the same coefficients. That expression is not the one the coefficients were fitted for. It gives the wrong element weighting: for silicon it returns $f(0) = 5.8$ where the correct value is $Z = 14$, and it decays too quickly with s . Both referees who inspected the code identified the discrepancy. The module was corrected to the expression above and re-validated (below). The production checkpoint (v21) was retrained from scratch with the corrected module and the corrected x_0 -mode loss of §3.4, under otherwise identical settings; the retrained checkpoint (v22) is the one behind Tables 2, 2a and 3 and Figs 4 and 5, and §5.2 and §5.7 give the effect. The exploratory Phase-4 rows of Table 1 were not retrained (§5.1).

Validation. On the first 1000 MP-20 test structures we compute the Pearson correlation between the differentiable pattern and the `pymatgen.XRD Calculator` reference, both on the 256-bin coarse grid. With the corrected form factor the mean is 0.988 and the median 0.999; 98.0 % of structures exceed 0.9. With the submitted version’s form factor the same structures give a mean of 0.952 and a median of 0.969; 90.7 % exceed 0.9. The correction improves the agreement for every one of the 1000 structures. On the first 50 structures the means are 0.939 (submitted form factor) and 0.974 (corrected). The submitted version reported 0.962 on 50 structures with the validation script’s default reflection range ($|h|, |k|, |l| \leq 10$), wider than the ≤ 5 training grid used above; Appendix C reproduces that value for the submitted expression (0.963 on the first 50, 0.960 on all 1000), gives 0.999 for the corrected expression at both sizes, and shows that the lower tail is a truncation effect of the ≤ 5 grid on small cells rather than a property of the form factor. Gradients through the structure factor with respect to fractional coordinates were verified non-zero.

As a loss. During training we recover an estimate $\hat{\mathbf{F}}$ of the clean coordinates from the model output (see §3.4), simulate $\hat{\mathbf{p}} = \text{DiffPXR}(\hat{\mathbf{F}}, \mathbf{Z}, \mathbf{L}_{\text{true}})$, and compare with the input pattern via a Pearson-correlation loss $\mathcal{L}_{\text{Debye}} = 1 - \rho(\hat{\mathbf{p}}, \mathbf{p})$. We use the true lattice for this auxiliary loss because the differentiable simulator is more sensitive to lattice errors than to coordinate errors at the early stages of training, and disentangling the two signals proved more stable.

3.4 Training objective

We use a VP-SDE with cosine schedule (Nichol & Dhariwal, 2021) for both channels:

$$\bar{\alpha}(t) = \cos^2\left(\frac{\pi}{2} \cdot \frac{t+s}{1+s}\right), \quad s = 0.008, \quad t \in [0, 1].$$

Coordinates are diffused on the flat torus $\mathbb{T}^3$ by wrapping the noisy sample to $[0, 1)^3$; the loss is computed on the periodic difference $(\hat{\mathbf{F}} - \mathbf{F} + 0.5) \bmod 1 - 0.5$. Lattice parameters are diffused in $\mathbb{R}^6$ after standardisation by the train-set mean and standard deviation.

The full training objective (best configuration, run `gpu_v13`) is

$$\mathcal{L} = \mathcal{L}_{\text{coord}} + \lambda_{\text{lat}} \mathcal{L}_{\text{lat}} + \lambda_{\text{aux}} \mathcal{L}_{\text{aux}} + \lambda_{\text{Debye}} \mathcal{L}_{\text{Debye}},$$

with $\lambda_{\text{lat}} = 0.1$, $\lambda_{\text{aux}} = 0.5$, $\lambda_{\text{Debye}} = 1.0$.




Figure 1: **Figure 1. Differentiable-simulator validation.** Distribution of Pearson correlation between the PyTorch DiffPXRd module and `pymatgen.XRD Calculator` over the first 1000 MP-20 test structures, with the corrected form factor (mean 0.988, median 0.999) and with the form factor of the submitted version (mean 0.952, median 0.969). The correction improves every structure. The residual lower tail is dominated by layered structures with strong texture in the reference.

Two non-trivial parameterisation choices. Both are essential and were arrived at by ablation, not foresight.

x_0 -residual prediction. Rather than asking the model to predict noise ϵ , the heads predict a **residual** that is added to the noisy input to obtain the clean estimate:

$$\hat{\mathbf{F}}^{(0)} = \mathbf{F}^{(t)} + \text{Coord-Head}(\mathbf{h}), \quad \hat{\ell}^{(0)} = \ell^{(t)} + \text{Lat-Head}(\cdot).$$

Because the heads are MLPs initialised to output near-zero, this means the model starts as the identity transform: at $t \approx 1$ (near pure noise) the model outputs the noisy input as its $\hat{\mathbf{F}}^{(0)}$ estimate, which is wrong but at least bounded; ϵ -prediction with the same architecture had to learn the entire transform from random initialisation and never recovered. At single seed this change appeared to lift the headline match metric $3.5\times$; a multi-seed checkpoint sweep (§7) shows that match-rate lift does not replicate, but the parameterisation does reproducibly improve the pattern-Pearson ($\approx 0.36 \rightarrow 0.40$, see §5.1).

Correction made during review. In x_0 -residual mode the Debye loss must be evaluated on the clean-coordinate estimate $\hat{\mathbf{F}}^{(0)} = \mathbf{F}^{(t)} + \text{Coord-Head}(\mathbf{h})$. The submitted code evaluated it on the residual $\text{Coord-Head}(\mathbf{h})$ alone. For every x_0 -residual run (v13, v14, v15, v16 and the Phase 9 checkpoint v21) the Debye term therefore compared the pattern of the wrong tensor with the target. The ϵ -mode runs v10 and v11 were unaffected, as were the coordinate and lattice losses and every evaluation-time number. The code is fixed. The production checkpoint was retrained with the corrected loss and the corrected simulator (v22, §3.3); its final EMA loss fell from 0.726 to 0.602, with the effect on the results in §5.2. The Phase-4 rows v13–v16 were not retrained and remain labelled exploratory single-seed runs carrying this defect (§5.1).

Lattice-head input. The lattice head reads the noisy lattice $\ell^{(t)}$ as an explicit input feature. Without

this, i.e. if the head only sees the (clean) lattice that is also passed to the denoiser as the geometry context, the head can only ever predict the unconditional mean noise, $\mathbb{E}[\epsilon] = 0$, and its loss stays pinned at the random baseline of 1.0 forever. With this fix, the lattice loss drops to 0.02–0.06 on stable runs.

We also experimented with two extensions that did not pan out (§5.4): a **Wyckoff-site embedding** added to atom features, and an auxiliary **distance-matrix loss** in which a pairwise MLP predicts ground-truth periodic distances between atoms. Both are documented in the released code behind the `--use-wyckoff` and `--dist-weight` flags.

3.5 Classical indexing as a control for the learned lattice head

The lattice head described in §3.2 is the *only* path by which the model encodes absolute d-spacings: every Bragg reflection 2θ position is set by the lattice through Bragg’s law, so the lattice parameters are the absolute reference frame for every peak position in the input pattern. As §5.5 shows, the encoder + denoiser learns *relative* d-spacing structure (Pearson 0.43 between predicted and target patterns at evaluation time, vs 0.97 between target and ground truth) but not the absolute scale. Classical autoindexing, which extracts peak positions, maps them to Q-space ($Q = 4\pi \sin\theta/\lambda$), and fits a unit cell by enumeration of $(h,k,l) \rightarrow Q^2(h,k,l)$, has solved this sub-problem since the 1970s and runs on a CPU in milliseconds per pattern.

We implement a from-scratch Q-space autoindexer that takes (i) the same simulated PXRD pattern fed to the encoder and (ii) the true crystal system of the target structure, supplied as a Bravais code, and returns a candidate $(a, b, c, \alpha, \beta, \gamma)$. Item (ii) is prior information beyond the composition that the learned head receives; it makes the indexer an optimistic control rather than a fair competitor (§7; the unknown-system run below removes it). Peak picking uses a 1D local-maximum filter with a relative-intensity floor; the candidate Q-vector is fit by least squares to the de-Wolff dichotomy parameter form for each Bravais lattice. At sampling time we run the indexer first, substitute its output for the lattice channel of the diffusion sampler, and run DDIM only on the coordinate channel.

Per-system accuracy (n = 1000 MP-20 test). Native indexer overall: 48.8 % strict cell match, 1.45 Å mean cell-length MAE. Per crystal system: hexagonal 77.9 % / 0.53 Å, orthorhombic 59.8 % / 0.96 Å, trigonal 43.3 % / 4.62 Å, monoclinic 18.9 % / 1.76 Å, triclinic 0 %. The low-symmetry blowup (monoclinic and triclinic each have ≥ 4 free cell parameters and the de-Wolff search becomes hypothesis-capped) is the dominant remaining error. A GSAS-II (Toby & Von Dreele, 2013) adapter (`src/pxrd_diff/indexer_gsas.py`) for the low-symmetry path ships behind a `--use-gsas` opt-in flag; it hangs on real MP-20 monoclinic/triclinic patterns inside `findBestCell` and is experimental (§7).

Without the crystal system. At a referee’s request the indexer was also run with no crystal system supplied. Cubic, tetragonal, hexagonal/trigonal, orthorhombic and monoclinic cells are searched in that order, highest symmetry first, and a lattice type is accepted by the rule the classical programs use (de Wolff, 1968; Werner *et al.*, 1985; Boulton & Louër, 2004): the first type whose best cell indexes every observed line with $M20 \geq 10$ is kept; if none does, the highest-M20 cell among those indexing every line; the reported system is the metric symmetry of the returned cell. On the same $n = 1\,000$ this reaches **43.4 %** strict cell match [40.4, 46.5] with a 1.67 Å mean cell-length MAE, against 48.8 % [45.7, 51.9] and 1.45 Å with the system given, and it recovers the lattice type itself for 65.4 % of the patterns (68.1 % of the 961 that both modes index). Per crystal system, strict success with the system given $\rightarrow$ unknown: cubic 52.9 $\rightarrow$ 49.2 %, hexagonal

77.9 $\rightarrow$ 65.4 %, trigonal 43.3 $\rightarrow$ 30.9 %, tetragonal 52.8 $\rightarrow$ 46.7 %, orthorhombic 59.8 $\rightarrow$ 53.6 %, monoclinic 18.9 $\rightarrow$ 19.6 %. Paired on the same structures, removing the system loses 65 cells and gains 9 (McNemar $p < 10^{-1}$): under ideal peak positions the crystal-system prior is worth about five percentage points to the indexer, not the difference between working and failing. A first implementation that simply kept the highest M20 across lattice types reached only 24.3 %, choosing the correct system for 40.7 % of patterns and cubic for 12 of the 238 cubic structures. M20 carries no penalty for free parameters, so on a pattern with a handful of exactly placed lines a four-parameter monoclinic sub-cell out-scores the true cubic cell (ScAlRh, mp-867922: M20 = 2 182 against 1 231, both indexing every line). That run is released beside the final one (`--unknown-rule maxm20`) because the comparison is itself informative: the choice of acceptance rule (24.3 $\rightarrow$ 43.4 %) matters more than the prior it replaces (43.4 $\rightarrow$ 48.8 %). Substituted into the sampler at three seeds, these cells recover 1.4 % [1.0, 1.9] (42 / 3 000) against 1.6 % [1.2, 2.1] with the system given: six structures lost and none gained on the same seeds (McNemar $p = 0.03$), with identical pattern Pearson (0.41), and no significant difference from the learned head’s 1.9 % (36 fixed, 52 broken, $p = 0.11$). Table 2 carries this as its third row.

The indexer is a drop-in in the mechanical sense: training is unchanged, and the diffusion sampler is unchanged except for the lattice-channel substitution at $t = T$. Patterns for which the indexer returns no cell (39 of the 1000 test patterns) are scored as misses; the submitted version had given them the true lattice (§5.7). §5.2 reports both with and without the indexer.

3.6 Sampling

DDIM (Song *et al.*, 2021) with 50 steps and $\eta = 0$. We start from independent standard Gaussian noise on the lattice and on the (un-wrapped) coordinates, then alternate the standard DDIM update on each channel. Coordinates are wrapped to $[0, 1)$ after every step. Predicted lattice parameters are de-standardised at the end and clipped to physically valid ranges $(a, b, c) \in [0.5, 100]$ Å, $(\alpha, \beta, \gamma) \in [10^\circ, 170^\circ]$ before being passed to `pymatgen.Lattice.from_parameters`.

A subtle bug fix is worth noting. In the x_0 -residual variant, recovering the implicit ε for the DDIM update requires dividing by $\sqrt{1 - \bar{\alpha}(t)}$; near $t = 0$ this denominator vanishes. We clamp it to 0.05 and skip the very last DDIM step when sampling, which removed a class of structures with NaN coordinates that we initially saw.

4. Experiments

Dataset. Canonical CDVAE (Xie *et al.*, 2022) MP-20 split: 27 136 / 9 047 / 9 046 train/val/test, ≤ 20 atoms/conventional cell. Patterns simulated with `pymatgen.XRD Calculator` (Cu $K\alpha_1$, $\lambda = 1.54184$ Å), max-normalised, 4 251 bins 5–90° / 0.02°; cached as `.npz`, zero failures over 45 196 structures.

Evaluation. Three views: (1) **composition match** (trivial, $\mathbf{Z}$ given); (2) **structure match** via `pymatgen.StructureMatcher` at $(\ell_{\text{tol}}, s_{\text{tol}}, \alpha_{\text{tol}}) = (0.2, 0.3, 5^\circ)$, the headline “match rate”; (3) **coordinate RMSD** on aligned, permuted atoms when matched. We also report space-group match at $\text{symprec} \in \{0.01, 0.05, 0.1, 0.2\}$, pattern Pearson, and R_{wp} .

The stricter “**all-correct**” rule combines `composition  $\wedge$  sg-match@symprec=0.1  $\wedge$  rmsd  $\leq$  0.1 Å`, the experimentally relevant case for downstream crystallographic use.

Two evaluation modes: *full pipeline* (predicted lattice + coords) and *true-lattice/coord-only* (ground-truth lattice substituted). The latter isolates coordinate quality from lattice quality and is the stricter ablation setting. In the indexer mode (§3.5) a pattern the indexer cannot index counts as a miss; it is never scored with the true lattice.

Implementation. AdamW, learning rate 10^{-3} (the value stored in the released v21 checkpoint; §5.7(d) on the submitted text’s 5×10^{-4}), cosine decay to zero over 100 k steps, WD 10^{-4} , grad-clip 1.0. Batch 64, ~24 epochs, single RTX 5090 (~1.5–1.7 h/run). 3.7 M parameters; $L = 5, d = 384$ (~10 M) gave no benefit (§5.2). VP-SDE with $t \sim U(0, 1)$. Total compute: ~35 GPU-hours, ~USD 25 on Vast.ai (Phase 4 ablation + Phase 9 v21 retrain + indexer/perturbation sweeps + DGpt n=1000 + PXRdnet n=20 + evaluation).

5. Results

Order: Phase 4 architectural ablation (§5.1) → Phase 9 indexer drop-in (§5.2) → reproduced baselines (§5.3) → failure catalogue (§5.4) → lattice-perturbation study (§5.5).

5.1 Phase 4: architectural ablation (true-lattice setting)

Table 1: seven runs sweeping the Debye loss, the lattice-input fix, x_0 -residual, and two extensions (Wyckoff, distance loss), all on $n = 1\,000$ MP-20 test with true lattice substituted. Every row is a single training seed. The table is exploratory: it guided the design, and only the ε -vs- x_0 -residual contrast was later replicated across seeds (§7). Runs v11 and v13–v16 were trained with the submitted version’s simulator, and v13–v16 with the x_0 -mode Debye defect (§3.3, §3.4). These rows were not retrained: only the production checkpoint was (v22, §5.2), and Table 1 is retained as the exploratory, single-seed record it always was, with the defects it carries stated here rather than corrected.

Table 1. Phase 4 coordinate-only ablation, MP-20 test ($n = 1\,000$, true lattice substituted). Single seed per row; exploratory.

Run	Parameterisation	λ_{Debye}	Wyckoff	λ_{dist}	Match %	Pearson	RMSD (Å)
v10	ε	0	—	0	1.40	0.359	0.17
v11	ε	1	—	0	0.90	0.365	0.15
v13	x_0-residual	1	—	0	2.51	0.434	0.22
v14	x_0 -residual	1	yes	0.01	0.80	0.367	0.14
v15	x_0 -residual	1	yes	0	2.10	0.392	0.21
v16	x_0 -residual	1	—	0.01	1.80	0.368	0.22

The **lattice-input fix** (§3.4) is load-bearing: it drops lattice loss from 1.0 (pinned at the prior variance for 100 k steps) to ~0.02 once $W_{\text{lat}}\ell^{(t)}$ enters the head input. A second, apparent lift did **not** survive replication: switching from ε to **x_0 -residual** at fixed Debye $\lambda = 1$ (v11 → v13) lifted single-seed match 0.9 % → 2.5 %, but a multi-seed checkpoint sweep (two seeds × five checkpoints, 20k–100k) found ε and x_0 -residual indistinguishable on match rate (§7), so the 3.5× match lift




Figure 2: **Figure 2. Phase 4 architectural ablation (single seed per run; exploratory).** Match rate (a) and pattern Pearson (b) for the seven coordinate-only runs of Table 1 on MP-20 test ($n = 1000$, true lattice substituted). The $\varepsilon \rightarrow x_0$ -residual switch (v11 $\rightarrow$ v13) shows the largest single-seed match lift (panel a), but that match lift does not replicate across seeds; only the pattern-Pearson gain (panel b) does (§7); adding Wyckoff embeddings or a distance-matrix loss regresses below it.

was a single-seed fluctuation. The x_0 -residual parameterisation does reproducibly improve pattern-Pearson ($0.36 \rightarrow 0.40$), the part of the Phase-4 signal that survives. Capacity is not the bottleneck: a 10.1 M-parameter run ($d = 384$, $L = 5$, `gpu_v12`) stayed at the same ~ 1.0 coord-loss plateau as v10. Under the full pipeline (sampled lattice), v13 drops to 1.2 % at $n = 256$; the Phase 9 checkpoint reaches 5.5 % under oracle (v22, three seeds; the submitted v21 gave 4.5 %) but only 1.9 % under the full pipeline. §5.2 attacks that gap.

5.2 Phase 9: classical indexer control

Table 2 compares four lattice sources on the same $n = 1\,000$ at three seeds with the v22 checkpoint (the v21 configuration retrained under the corrected simulator and x_0 -mode Debye loss, §3.3, §3.4, §5.7); the rows differ only in the lattice the sampler conditions on. The indexer rows are controls with extra prior information (§3.5). The submitted version’s v21 numbers are retained in §5.7.

Table 2. Phase 9 lattice-source comparison ($n = 1\,000$ MP-20 test $\times$ three seeds, 3 000 structure-seeds pooled, v22 checkpoint). The indexer receives the true crystal system; the 117 structure-seeds it cannot index are scored as misses. Match % with 95 % Wilson CI; the other columns are means over seeds.

Lattice source	Match % (95 % CI)	All-correct %	sg@0.1 %	RMSD med (Å)	R_wp	Pearson
Learned head only (full pipeline)	1.9 [1.5, 2.5]	0.0	1.4	0.232	17.0	0.18
Classical Q-space autoindexer (§3.5), crystal system given	1.6 [1.2, 2.1]	0.2	0.3	0.035	11.4	0.41

Lattice source	Match % (95 % CI)	All- correct %	sg@0.1 %	RMSD med (Å)	R_wp	Pearson
Classical Q-space autoindexer, crystal system unknown	1.4 [1.0, 1.9]	0.2	1.6	0.038	12.2	0.41
Ground-truth lattice (oracle)	5.5 [4.7, 6.4]	0.8	2.2	0.044	5.1	0.71

A three-seed paired test. Because both arms see the same materials, a paired McNemar test on per-structure outcomes, not overlapping Wilson intervals, decides whether two lattice sources differ. Every source was run at three seeds ($n = 1\,000$ each), differing only in the lattice channel; with v22 every arm scores all 3 000 structure-seeds (the submitted v21 checkpoint sampled degenerate cells that hung **StructureMatcher**/**spglib**, leaving 1 949 scorable learned-head rows). The indexer recovers **42** structures the learned head misses and the learned head recovers **52** the indexer misses (Table 2a).

Table 2a. Paired per-structure outcomes, learned head vs indexer control (v22, three seeds pooled; match = **StructureMatcher** hit).

	indexer match	indexer miss	row total
learned match	6	52	58
learned miss	42	2 900	2 942
col total	48	2 952	3 000

The discordant cells are $b = 42$ (indexer fixes), $c = 52$ (indexer breaks); a two-sided McNemar test gives $\mathbf{p} = \mathbf{0.35}$: the two lattice sources tie on match rate. The submitted version reported the opposite: with v21 the learned head matched 0 of 1 949 scored structure-seeds against the indexer’s 46 of 2 883, 27 fixed and none broken ($p = 1.5 \times 10^{-8}$). That lead was a property of the defective checkpoint (§3.4): paired on the same seeds, the retrain gains 32 learned-head structures and loses none ($p = 4.7 \times 10^{-1}$), while the indexer arm moves by 12 gained and 10 lost ($p = 0.83$). The claim that the control improves on the learned head is withdrawn as a general statement. What survives is the per-system pattern: the indexer still fixes hexagonal (16 fixed, 0 broken) and tetragonal (7, 2) structures, the systems where its cell error sits below the 0.5 Å knee (§5.5), but the corrected learned head wins the cubic cases (18 fixed, 47 broken; learned 52 / 714, indexer 23 / 714, oracle 96 / 714). The indexer’s cells remain far closer to the truth where the match rates tie (pattern Pearson 0.41 against 0.18; RMSD median 0.035 against 0.232 Å), so in the cubic cases the match rate saturates on coordinate decoding rather than on the cell. On the stricter all-correct metric the indexer fixes 6 and breaks 0 ($p = 0.03$).

What neither lattice source does is close the gap to the oracle: **5.5 % [4.7, 6.4]** at three seeds (165 / 3 000; 4.5 % [3.8, 5.3] with the submitted checkpoint, unpaired $p = 0.08$). Paired against the learned head it fixes 134 structures and breaks 27 ($p = 3 \times 10^{-18}$), and against the indexer the interval is disjoint. This is the effect that replicates across checkpoints and seeds: even with a perfect lattice, match caps at 5.5 %, so coordinate prediction is the dominant limiter while lattice error explains the $1.9 \rightarrow 5.5$ % shortfall. The indexer’s per-system MAE (1.45 Å overall, dominated by mono 1.76 Å and triclinic NaN) sits above the 0.5 Å knee; closing the rest requires fixing the low-symmetry path (GSAS-II remains experimental).




Figure 3. Per-system indexer performance. (a) Strict indexing rate with the crystal system given (blue) and with the system unknown (orange; acceptance rule of §3.5), and the consistent rate with the system given (light blue), on $n = 1\,000$; dashed line, overall strict rate given / unknown. (b) Length MAE with the system given. High-symmetry systems (hex, ortho) sit below the 0.5 Å knee; trig and mono blow up; triclinic is unindexable in this implementation. Removing the crystal system costs a few points in every system except monoclinic and never changes the regime. The overall 1.45 Å MAE hides two regimes, and explains why the indexer’s paired gains over the learned head (§5.2) are confined to the hexagonal and tetragonal systems.

5.3 Head-to-head with external baselines

DiffractionGPT (knc6/diffractgpt_mistral_chemical_formula) at $n = 1\,000$, PXRNet (therealgabeguo/cdvae_xrd_sinc100) at $n = 20$ (limited by ~ 4 GPU-hr/material), and deCIFer (deCIFer_v1, conditioned on XRD + composition) at $n = 300$ re-scored on *our* StructureMatcher harness at matching tolerances. Crystallize’s checkpoint download is inactive; cite-but-not-reproduced.

Table 3. Reproduced head-to-head on MP-20 test (StructureMatcher ltol 0.2, stol 0.3, angle_tol 5°). The first PXRNet-Diff row is the learned head, which receives the pattern and composition only, the same information as the external models; the second is the classical-indexer control, which is additionally supplied with the true crystal system. Both are the v22 checkpoint at three seeds (§5.2). The rows set the scale of the diagnostic model and are not a competitiveness claim.

Model	n	Match % (95 % CI)	All-correct % (95 % CI)	sg@0.1 %	RMSD med (Å)
PXRNet-Diff (v22, learned head; three seeds pooled)	3 000	1.9 [1.5, 2.5]	0.0 [0.0, 0.1]	1.4	0.232
PXRNet-Diff + indexer control, crystal system given (v22; three seeds pooled)	3 000	1.6 [1.2, 2.1]	0.2 [0.1, 0.4]	0.3	0.035
DiffractionGPT (Choudhary, 2025)	1 000	18.9 [16.6, 21.4]	15.9 [13.8, 18.3]	21.3	0.001
PXRNet sinc100 (Guo <i>et al.</i> , 2025)	20	30.0 [14.5, 51.9]	5.0 [0.9, 23.6]	5.0	0.011
deCIFer (Johansen <i>et al.</i> , 2025)	298	73.8 [68.6, 78.5]	69.5 [64.0, 74.4]	83.6	0.009

Match rate ranks deCIFer $\gg$ PXRNet $>$ DGpt $>$ PXRNet-Diff. deCIFer, a concurrent autoregressive CIF-token model conditioned on XRD + composition, is the strongest external reference at 73.8 % match / 69.5 % all-correct ($n = 300$); its all-correct nearly equals its match, the same CIF-token signature seen in DGpt. PXRNet’s lead over DGpt is consistent with its 35 k-op latent optimisation, whereas DGpt is one forward pass. PXRNet-Diff, at 1.9 %, sits an order of magnitude below all three, as expected for a 3.7 M-parameter diagnostic model. The *all-correct* ordering need not follow match rate: DGpt’s all-correct (15.9 % [13.8, 18.3]) sits close to its match, while PXRNet’s drops from 30 % match to 5 % all-correct, i.e. PXRNet matches structure under loose tolerances but rarely recovers the space group, whereas DGpt’s CIF-token output appears to encode symmetry. **We flag this as a hypothesis, not a finding.** At $n = 20$ the PXRNet all-correct estimate (1/20) carries a 95 % CI of [0.9, 23.6], which overlaps DGpt’s [13.8, 18.3]; confirming the reversal would require $n \approx 100\text{--}200$. §6 says why it would matter. **RMSD** is heavy-tailed for us and tight-near-zero for both baselines, suggesting our matches are tolerance-stretch rather than coordinate recovery.

What the gap does and does not reflect. All four systems are given the composition (§3.1). The learned-head row of PXRNet-Diff receives nothing else; the indexer-control row additionally gives its indexer the true crystal system (§3.5), so on prior information that row is favoured. On the pattern itself the asymmetry runs the other way: DiffractionGPT is conditioned on a **coarse 300-bin pattern** ($2\theta \in [0, 90^\circ]$, 0.3° bins, rendered as a ;-separated string), an order of magnitude lower

resolution than PXRD-Diff’s 4 251-bin 0.02° input, yet it still reaches 18.9 %, so the external models do not win by seeing more of the pattern. The differences that matter are **capacity** (DiffractionGPT is a 7 B-parameter LLM; PXRD-Diff is 3.7 M, three orders of magnitude) and **inference compute** (PXRDnet’s ~ 35 k-op latent optimisation vs our single DDIM pass). Table 3 shows the expected gap between a deliberately small diagnostic model and production-scale systems; it sets the scale and does not rank PXRD-Diff among them (§6).




Figure 4: **Figure 4. Three-way headline comparison.** Match rate, all-correct, and pattern-Pearson on the same MP-20 test materials scored through *our* StructureMatcher harness, with 95 % Wilson CIs. PXRD-Diff is the corrected v22 checkpoint with its learned lattice head, no prior beyond composition, three seeds pooled ($n = 3 \times 1\,000$); its indexer control reaches 1.6 % and its true-lattice oracle 5.5 % (§5.2). It sits an order of magnitude below the external systems on match rate and above them on per-pattern Pearson (0.18 against 0.02 and 0.16). Note the wide PXRDnet ($n = 20$) intervals: the DGpt > PXRDnet all-correct ordering is a hypothesis, not an established result (§5.3, §6).

Caveats. PXRDnet’s $n = 20$ intervals (roughly ± 20 pp) and the preprocessing differences between the three pipelines (identical material_ids, pattern inputs not bit-identical) are bounded in §7.

5.4 What did *not* work

Four interventions failed at $n \geq 200$; one paragraph each.

Top-K Debye rerank (Phase 9.1.4). Top-5 indexer candidates reranked by Debye loss: match 1.6 % $\rightarrow$ 1.0 %, all-correct 0.1 % $\rightarrow$ 0.0 % at $n = 1\,000$. The indexer surfaces more wrong cells faster than rerank can filter.

Debye-gradient guidance during DDIM (Phase 9.2). Coordinate-channel guidance term $-g \cdot \nabla_{\mathcal{L}_{\text{Debye}}}$, sweep $g \in \{0, 0.5, 1, 2, 5\}$ on $n = 200$: flat 1.0–2.5 % within noise, all-correct 0 % throughout. Sampling-time Debye gradient is too noisy to steer the trajectory, consistent with Segal *et al.* (2026), who show the powder-XRD similarity loss landscape is too non-convex for direct gradient descent.

Wyckoff-letter embeddings (v15, $n = 1\,000$). Each atom received an `nn.Embedding(27, 256)` token for its `spglib`-assigned Wyckoff letter. A Wyckoff letter is meaningful only within a

given space group: the same letter denotes different multiplicities, site symmetries and coordinate constraints in different groups (Dauter & Jaskolski, 2010; de la Flor *et al.*, 2023), so the embedding carried a label, not a constraint, and nothing in the loss or the sampler enforced the fixed or symmetry-related coordinates a Wyckoff position implies. The embedding learns (5 404 of 6 912 entries non-zero), but match drops from 2.5 % to 2.1 % and lattice prediction destabilises (lattice loss ~ 0.6 vs ~ 0.02 ; Figure 5), presumably because the inflated atom-feature norm shifts the lattice-pool input distribution. The negative result is narrow: a space-group-agnostic label token does not help and may hurt training dynamics; whether symmetry imposed as a constraint, on the coordinates or on the space group, would help was not tested.




Figure 5: **Figure 5. Wyckoff embedding destabilises lattice prediction.** Lattice (a) and coordinate (b) training-loss curves for the distance-aux run (v16) vs the Wyckoff-embedding run (v15). Adding the Wyckoff embedding inflates the lattice loss by an order of magnitude (~ 0.6 vs ~ 0.02) while leaving the coordinate loss largely unchanged; the failure is localised to the lattice channel.

Distance-matrix aux loss (v16, $n = 1\,000$). Pairwise MLP predicting periodic distances, $\lambda_{\text{dist}} = 0.01$ (largest stable). Loss decreases (final ~ 2.0 vs ~ 10 random) but match drops 2.5 % $\rightarrow$ 1.8 %; biases features toward absolute-distance reconstruction rather than the relative updates diffusion needs.

Combination (v14). Wyckoff + distance loss together collapses to 0.8 %, below the bare ε baseline; the two interact destructively, no clean theoretical explanation.

5.5 Lattice-perturbation study

Best predicted-vs-target Pearson is 0.43; the same model with *correct* coordinates plugged into the differentiable simulator hits 0.97; the encoder + denoiser leaves ~ 0.5 of pattern-space agreement on the table. The aux head (sees **g** only) reaches a 0.007 *normalised* lattice-regression loss, but that figure is a relative fit that does **not** survive conversion to absolute scale: when the aux head’s own lattice is substituted into the sampler it recovers only 1.2 % (§5.6), because its 1.1 Å length MAE shifts every Bragg reflection by ~ 15 %. So the lattice head delivers crystal-system and relative-spacing structure but not absolute d-spacings at the <0.5 Å precision the knee demands; §5.6

shows that giving the head the peak positions does not repair this.

Perturbing the sampler-supplied lattice by Å of cell-length error, v21 at $n = 200$: match decays 5.6 % ($\Delta=0$) $\rightarrow$ 2.0 % (0.5 Å) $\rightarrow$ 0.5 % (1.0 Å) $\rightarrow$ 0 % (≥ 1.5 Å), steeper than linear with a knee at ~ 0.5 Å. This matches the indexer’s per-system MAE: high-symmetry systems (hex, ortho; $\text{MAE} \leq 1.0$ Å) lift; low-symmetry (mono, tri; $\text{MAE} \geq 1.5$ Å) do not. §5.2’s per-system indexer gains are *predictable*: they occur where the cell error is below the knee.

5.6 Why the learned lattice fails: neither the head nor the pooling repairs the scale

The natural reading of the 0.007 aux loss, that the lattice is already recoverable from the encoder and only the denoiser is at fault, is testable, and it is wrong. Two experiments (released as `paper/phase5_results/` and `paper/phase5b6_results/`) promote the auxiliary head from a training-time regulariser to the inference-time lattice source and measure what it actually delivers.

Aux head as lattice predictor (no retrain). Substituting the v13 aux head’s predicted $(a, b, c, \alpha, \beta, \gamma)$ for the sampled lattice channel at $n = 1\,000$ gives a length MAE of (1.12, 1.02, 1.38) Å and an angle MAE of (13.0°, 11.9°, 17.6°) at **99.8 %** cell validity, far better conditioned than the diffusion sampler, whose predicted cells are near-0 % valid, but still ~ 15 % relative length error. End-to-end this recovers **1.2 %** match (1.3 % with Phase 4 ensembling + coordinate refinement) and 0.0 % all-correct, with pattern Pearson collapsing to **0.013**: a 15 % length error shifts every Bragg reflection by a similar fraction in 2θ , so the simulated pattern no longer overlaps the target.

Stronger, constrained supervision (retrain). Bounding the head (softplus lengths, sigmoid-scaled angles), raising the aux weight $0.5 \rightarrow 5.0$, and warm-starting 30 k further steps leaves the MAE essentially unchanged ($a \approx 1.11$ Å, $\alpha \approx 14^\circ$). No head reading **g** recovers the absolute scale. The submitted version read this as a property of the *pooled global encoding*, on the grounds that the `AdaptiveAvgPool1d(1)` collapse §3.2 identified for coordinates would also discard the peak-position fidelity Bragg’s-law decoding needs; the controlled ablation below tests that reading and does not support it.

This reconciles the diagnosis with the indexer result. The 0.007 normalised aux loss is a *relative* fit, enough to place the crystal system (a space-group head on the same encoder reaches 39 % top-1 / 71 % top-5 across 230 classes, released under `phase5b6_results/`) but not the absolute scale. The autoindexer succeeds where every learned head fails because it assigns Miller indices to peak *positions* by explicit search rather than regressing six numbers from a summary, which is why its paired gains in §5.2 lie in the high-symmetry systems where peak positions over-determine the cell, and why head-side fixes (Phase 5 here, top-K rerank in §5.4) cannot substitute for it. The ablation below asks whether it is the summary, rather than the regression, that fails.

Controlled pooling ablation. The two experiments above vary the head and its supervision with the encoder fixed; they do not by themselves isolate global average pooling as the cause, since the encoder, the information content of its representation, the head parameterisation and optimisation difficulty remain confounded, as one referee noted. To separate them, a controlled ablation holds the encoder, the denoiser, the training objective and the schedule fixed and changes only how the auxiliary head reads the encoder: (i) the global average pool **g** of the submitted version; (ii) a position-aware pooling that retains the 2θ coordinate of each feature; (iii) explicit peak-position features extracted from the input pattern and concatenated to **g**. If absolute d-spacings are lost at

the pooling stage, variants (ii) and (iii) should reduce the ~ 1.1 Å length error while (i) does not. Table 4 gives the result on both encoders.

Table 4. Controlled pooling ablation ($n = 1\,000$ MP-20 test). Three heads trained under an identical loss, AdamW 10^{-3} , batch 64, sigmoid-bounded outputs, on a frozen encoder: the submitted v21 encoder at three epochs and the retrained v22 encoder (§5.2) at ten. Length and angle MAE are means over (a, b, c) and (α, β, γ) ; “ < 0.5 Å” is the fraction of structures whose mean length error lies below the knee of §5.5; the last column substitutes the head’s cells into the sampler at three seeds ($n = 3\,000$ structure-seeds; the full pipeline reaches 1.9 % and the indexer 1.6 %, Table 2). The checkpoint rows are the head trained jointly for 100 k steps, without retraining.

Head	Params	Encoder, epochs	Length MAE / Å	Angle MAE / °	Median length error / Å	< 0.5 Å	End-to-end match % [95 % CI]
Checkpoint – aux head	–	v21, –	1.32	15.6	0.90	0.33	–
(i) global average pool g	67 k	v21, 3	1.29	16.0	0.87	0.29	–
(ii) position-aware attention pooling	531 k	v21, 3	1.31	16.4	0.94	0.24	–
(iii) explicit peak features + g	209 k	v21, 3	1.28	16.1	0.88	0.28	–
Checkpoint – aux head	–	v22, –	1.28	15.0	0.90	0.31	–
(i) global average pool g	67 k	v22, 10	1.29	15.5	0.91	0.28	0.7 [0.5, 1.1]
(ii) position-aware attention pooling	531 k	v22, 10	1.21	15.4	0.88	0.29	0.5 [0.3, 0.9]
(ii) position-aware attention pooling	531 k	v22, 30	1.12	14.6	0.75	0.36	0.8 [0.5, 1.2]

Head	Params	Encoder, epochs	Length MAE / Å	Angle MAE / °	Median length error / Å	< 0.5 Å	End-to- end match % [95 % CI]
(iii) explicit peak features + g	209 k	v22, 10	1.28	15.4	0.90	0.30	1.2 [0.9, 1.7]

The prediction is not borne out on either encoder. On v21 the three arms lie within 0.03 Å of one another (~2 % of the error). On v22, with ten epochs, arms (i) and (iii) converge by the sixth epoch to 1.29 and 1.28 Å, and arm (ii) reaches 1.21 Å at ten epochs and 1.12 Å at thirty, where it is converged (flat over the last six epochs), a 13 % reduction. Every arm sits more than 2.2× above the 0.5 Å knee, and arm (i) reproduces the checkpoint head on both encoders, so the budget suffices to match the baseline. End to end, the substituted cells recover 0.7 %, 0.5 % and 1.2 % for (i), (ii) and (iii), each below the 1.9 % of the full pipeline (paired McNemar $p = 7 \times 10^{-3}$, 3×10^{-3} and 8×10^{-3}) and the indexer’s 1.6 %, with pattern Pearson 0.08 in every arm; the head-level gain of arm (ii) does not carry through: its converged thirty-epoch cells recover 0.8 % [0.5, 1.2] (24 / 3 000, 22 of them cubic), level with arm (i). Every match obtained through a regression head is a cubic structure (18, 16 and 34 of the 22, 16 and 37), the one-parameter case; none is hexagonal or tetragonal, where the indexer’s search succeeds (§5.2). Restoring absolute peak positions to the head does not repair the absolute scale, and the submitted reading that global pooling discards the d-spacing information is withdrawn. Arm (iii) is decisive: it receives the same peak positions the autoindexer reads, yet a head trained to regress six cell parameters does not turn them into a cell, because that requires assigning Miller indices to peaks, a combinatorial step the autoindexer performs by explicit search and a pointwise-loss regressor does not learn here. This revised reading is the best-supported one, not a demonstrated cause; it accounts for every head-side failure (the two experiments above, the top-K rerank of §5.4), for the space-group head’s success on relative features, and for the indexer succeeding exactly where the assignment is over-determined (§5.2). The next gain must come from obtaining the cell by search, or from a model that performs the peak-to-index assignment explicitly, not from a different pooling, a better head or a better denoiser. Two caveats remain: the encoder is frozen, so an encoder trained from scratch under position-aware pooling could behave differently (not run); and the heads are trained on the frozen encoder’s features rather than jointly, so a jointly trained position-aware head is untested. Runs: `paper/phase15_results/pooling_ablation_v21/`, `pooling_ablation_v22/`, `v22_pool_*`.

5.7 Code audit and corrections

Review of the released code, prompted by two referees, identified three implementation errors. Each is listed with its effect on the numbers in this paper.

- (a) *Atomic form factor*. The differentiable simulator evaluated $\sum_k a_k \exp(-b_k s^2)$ with coefficients fitted for $Z - 41.78214 s^2 \sum_k a_k \exp(-b_k s^2)$ (§3.3). Affected: the Debye training loss in every run with $\lambda_{\text{Debye}} > 0$ (v11, v13–v16, v21); the Debye-gradient guidance and top-K rerank of §5.4; and the validation figure. Not affected: all match, all-correct, space-group and RMSD numbers, which `pymatgen` scores on the predicted structures; the indexer; the oracle;

the perturbation study; the aux-head substitution of §5.6; and the external baselines. The module was corrected and re-validated (mean Pearson $0.952 \rightarrow 0.988$ over 1000 structures). The production checkpoint was retrained from scratch with the corrected module and the fix of (b), under otherwise identical settings (v22; final EMA loss 0.602 against 0.726); Tables 2, 2a and 3 and Figs 4 and 5 now report it at three seeds, and §5.2 gives the paired before/after accounting (learned head $0 / 1\,949 \rightarrow 58 / 3\,000$, indexer $46 / 2\,883 \rightarrow 48 / 3\,000$, oracle $135 \rightarrow 165 / 3\,000$; the indexer’s lead became a tie). The Phase-4 rows of Table 1 and the §5.4 experiments were not rerun (~ 18 GPU-hours for the six configurations at three seeds); they keep their exploratory labels.

- (b) *Debye loss in x_0 -residual mode.* The loss was evaluated on the residual tensor rather than on the clean-coordinate estimate (§3.4). Affected: the Debye term of runs v13–v16 and v21. Not affected: the ε -mode runs v10 and v11, and every evaluation-time number. Fixed; the retrained runs of (a) also carry this fix.
- (c) *Indexer fallback.* When the classical indexer returned no cell, the sampling script substituted the true lattice and the pattern was scored as if indexed. This affected 39 of the 1000 test patterns in every indexer run, 117 structure-seeds over three seeds. Recomputing from the released per-structure records, none of the 117 fallback structure-seeds matched, so the submitted McNemar cells (27 fixed, 0 broken) and the all-correct cells (4, 0) were unchanged by the fallback; the submitted match rate becomes $46 / 2\,883 = 1.6\%$ [1.2, 2.1] on the indexed structure-seeds, in place of $46 / 3\,000 = 1.5\%$ [1.2, 2.0]. (Those v21 cells are superseded by the v22 retrain of (a), which is what §5.2 now reports.) The script now scores unindexed patterns as misses and records their count. The top-K rerank experiment of §5.4 used the same fallback; because the fallback rows did not match, its 1.0 % result can only be lower, and its conclusion (no improvement) stands.
- (d) *Reported learning rate.* The submitted §4 gave the learning rate as 5×10^{-4} . The arguments stored in the released v21 checkpoint give 10^{-3} , which is the value that was used; §4 now says so. No result depends on this, but the referees asked for a systematic check of the description against the implementation and this is what it found.

The corrected code, the per-structure records used for (c), the re-validation script and a regression test that pins the form-factor convention (`tests/test_debye_form_factor.py`) are in the released repository.

6. Discussion

Lattice-recovery bottleneck (the robust result). A 0.007 *normalised* aux loss shows **g** carries lattice-*relevant* structure, but it is a relative fit: feeding the aux head’s own lattice to the sampler recovers only 1.2 % (§5.6), and neither stronger constrained supervision nor restoring absolute peak positions to the head moves its ~ 1.1 Å error, so the pooling reading of the submitted version is withdrawn in favour of the regression formulation (§5.6). The paired oracle-vs-learned-head test (134 fixed, 27 broken, $p = 3 \times 10^{-18}$; §5.2) is the paper’s statistically firmest claim and the one result that replicates across the submitted and the corrected checkpoint. The Q-space autoindexer, which reads peak positions directly, recovers part of the 5.5 % oracle ceiling on the hexagonal and tetragonal subset; with the corrected checkpoint it no longer leads the learned head in aggregate, but its cells stay far closer to the truth (Pearson 0.41 against 0.18). Top-K rerank and gradient

guidance fail because they operate *after* the encoder commits to a bad lattice prior, while the indexer bypasses that decision. The oracle ceiling is itself only 5.5 %, so even a perfect lattice leaves coordinates as a co-limiter (§8).

Per-atom anchoring is hard. PXRD is permutation-invariant; our denoiser is permutation-equivariant; there is no symmetry-breaking signal that pins atom i to a specific Wyckoff site. The Wyckoff-letter embedding tried to break symmetry at the input with a label rather than a constraint, and failed (§5.4). A promising untried direction: break symmetry at the *output*: predict an unordered set of orbits plus a Hungarian-style matcher.

A measurement hypothesis: match rate may conflate two capabilities. PXRDnet’s 30 % match / 5 % all-correct vs DGpt’s 18.9 % / 15.9 % (Table 3) suggests the two systems recover *structure* at broadly comparable rates while differing sharply in *symmetry* recovery: DGpt’s CIF-token output appears to encode space group in a way PXRDnet’s coordinate decode does not. The numbers cannot yet support this (PXRDnet’s all-correct CI [0.9, 23.6] at $n = 20$ overlaps DGpt’s), so we raise it as a *measurement hypothesis* worth testing at adequate n . If true, a single “match rate” reported under inconsistent tolerances would conflate structure recovery with symmetry recovery, and downstream crystallography, where the space group matters as much as the coordinates, would be mis-served by it. The contribution here is the shared harness that makes such a test possible, not the underpowered comparison itself.

Niche for small reproducible models. At 1.9 % PXRD-Diff sits an order of magnitude below much larger systems, and that is the point of its size. Small models make the lattice-recovery bottleneck visible (capacity hides it) and make the indexer control trivial to wire in (larger models would need architecture-level surgery). The recipe (encoder + Phase 4 fixes + Phase 9 indexer control) is 3.7 M parameters and runs inference in seconds.

7. Limitations

- **Scope.** Composition given; simulated PXRD only (no instrument response, preferred orientation, asymmetry, background); MP-20 only, with no larger cells, organics, or higher-Z.
- **Statistics.** The Phase 4 ablations in Table 1 are single-seed. A multi-seed checkpoint sweep of the ε -vs- x_0 -residual contrast (two seeds $\times$ five checkpoints, 20k–100k, true-lattice $n = 1000$) shows the reported $3.5\times$ *match-rate* lift does **not** replicate: pooled match is 1.20 % (ε) vs 1.35 % (x_0 -residual), and x_0 never exceeds 1.5 % at any checkpoint, including the 80 k checkpoint that matches the original v13’s 79.5 k, so the single-seed 2.51 % was a high draw, not a checkpoint-selection effect. We therefore demote the *match-rate* lift to a single-seed artifact. What does reproduce is the pattern-Pearson improvement (ε 0.359 vs x_0 0.403, with disjoint ranges across both seeds), consistent with Table 1. The lattice-input fix remains load-bearing as a training-dynamics result (lattice loss $1.0 \rightarrow 0.02$), and the oracle-vs-learned-head gap is the effect that replicates (§5.2). The indexer-vs-learned-head comparison is the cautionary case: significant with the submitted checkpoint (27 fixed, 0 broken), a tie with the corrected one (42, 52; $p = 0.35$), so a paired test at three seeds protects against seed noise but not against a training defect that moves one arm (§5.2, §5.7).
- **Baselines.** PXRDnet $n = 20$ has a 95 % Wilson CI of $\approx \pm 20$ pp; the DGpt > PXRDnet all-correct ordering (§5.3/§6) is therefore a hypothesis, not a result. $n = 200$ needs ~ 33 days RTX 5090. Crystallize checkpoint download is inactive (verified 2026-06-01); cited but

unreproduced. Patterns are not bit-identical across the three preprocessors (PXRDiff 4 251-bin 2θ , DGpt 300-bin 2θ , PXRNet 4 096-bin Q with sinc^2 broadening): structures match, patterns do not; we did not quantify the residual preprocessing effect on match rate, and a matched-vs-native re-scoring on a structure subset is the natural check.

- **Indexer realism.** Two assumptions make the indexer’s accuracy an optimistic ceiling. (i) *Crystal system.* The main runs supply the true Bravais code to the de-Wolff fit (§3.5); the unknown-system run removes it at a cost of about five percentage points at the indexing stage ($48.8 \rightarrow 43.4$ % strict) and six of 48 downstream matches (1.4 % against 1.6 %), so the assumption changes little at either stage. (ii) *Patterns are ideal.* All patterns are clean `pymatgen` simulations with no zero-point shift, sample-displacement error, or peak asymmetry, precisely the perturbations that drive real-world indexing failure. The 48.8 % native-indexer match and the downstream 1.6 % should be read as upper bounds under ideal peak positions. A zero-shift/displacement robustness sweep ($n = 300$) confirms the concern: the indexer’s length-MAE rises from 1.24 Å on clean patterns past the v20 learned head’s 1.37 Å once displacement or zero-shift exceeds $\sim 0.10^\circ$, so the indexer’s cell-accuracy advantage over the learned head (§5.2) holds only for well-calibrated input.
- **Idealised conditions and the bottleneck conclusion.** Every pattern is an ideal simulation. Instrument profile, zero shift, sample displacement, preferred orientation and background all move or reshape peaks. Each would further degrade an encoder that already fails to fix absolute d-spacings on clean input. The lattice-recovery bottleneck diagnosis is therefore a lower bound on the difficulty of real data, not an estimate of it. The 0.5 Å knee is a property of the sampler under clean cells; under real peak shifts its position may move, and the indexer’s own degradation past $\sim 0.10^\circ$ suggests the whole curve shifts towards smaller tolerable errors.
- **Indexer (low-symmetry path).** GSAS-II low-symmetry path hangs in `findBestCell` on real MP-20 mono/tri patterns; shipped behind `--use-gsas` but experimental.
- **Training.** Debye loss uses ground-truth lattice; a curriculum gradually replacing true with predicted was not tried.

8. Conclusion

A 3.7 M-parameter conditional diffusion model, used as a diagnostic instrument, locates the dominant accessible failure in PXRDiff inversion in the recovery of the unit cell from the pattern, upstream of the denoiser. The encoder’s pooled representation supports a 0.007 normalised lattice-regression loss yet delivers a ~ 1.1 Å absolute cell error that no head reading it reduces, whether constrained, position-aware or handed explicit peak positions; the pooling reading is withdrawn and the evidence points to the regression formulation, which leaves the peak-to-index assignment unsolved. A true-lattice oracle caps match at 5.5 % [4.7, 6.4] (three seeds) against 1.9 % [1.5, 2.5] for the learned head (paired McNemar $p = 3 \times 10^{-18}$), so coordinate prediction remains a co-limiter. A classical Q -space autoindexer, supplied with the true crystal system, serves as a control: it reaches 1.6 % [1.2, 2.1] and, with the corrected checkpoint, ties the learned head at three seeds (42 fixed, 52 broken, $p = 0.35$), its gains confined to the hexagonal and tetragonal systems below the 0.5 Å knee. Re-scored on one `StructureMatcher` harness, deCIFer reaches 73.8 % / 69.5 %, DGpt 18.9 % / 15.9 % and PXRNet 30.0 % / 5.0 % (match / all-correct); these set the scale of the diagnostic model and are not a ranking claim, and the apparent divergence in their space-group recovery is a measurement hypothesis the $n = 20$ PXRNet sample cannot yet confirm. A code audit corrected

three implementation errors; retraining the production checkpoint under the corrected loss raised the learned head from 0 to 1.9 % and removed the indexer’s aggregate lead of the submitted version, leaving the oracle gap as the result that replicates. The lattice-input fix is load-bearing; the x_0 -residual switch improves pattern-Pearson but its single-seed match-rate lift did not replicate (§7); the remaining interventions are documented as failures. Code, checkpoints, all per-phase JSONs, per-structure records, both indexer paths and the differentiable Bragg module are released.

Acknowledgments

Use of AI tools. The author used a large language model (Claude, Anthropic) as a coding and writing assistant throughout this work, for drafting and editing text and code under the author’s direction. The author reviewed all code, analyses and text, and takes full responsibility for their content.

We thank the maintainers of `pymatgen`, `spglib`, the CDVAE benchmark, GSAS-II, and the upstream maintainers of DiffractionGPT (`atomgptlab/atomgpt`) and PXRNet (`gabeguo/cdvaexrd`) for releasing checkpoints and code that made the head-to-head reproduction in §5.3 possible. Compute was rented from Vast.ai; total spend was approximately USD 25 across roughly 30 GPU-hours on RTX 5090 instances (Phase 4 ablation, Phase 9 retrain + indexer sweeps, DiffractionGPT $n = 1000$ inference, and PXRNet $n = 20$ inference).

Author Contributions (CRedit)

F. Cai: Conceptualization, Methodology, Software, Validation, Formal analysis, Investigation, Data curation, Writing – Original Draft, Writing – Review & Editing, Visualization, Project administration.

Conflict of Interest

The author declares no competing interests.

Funding

This research received no external funding.

Data and Code Availability

Source code, trained checkpoints, and all per-phase training/evaluation logs are openly available in the GitHub repository at <https://github.com/fronkt/pxr-diff> and archived at Zenodo: <https://doi.org/10.5281/zenodo.20738994> (DOI: 10.5281/zenodo.20738994). The MP-20 dataset is publicly available via the CDVAE benchmark. All experiments reproduce from a single `requirements.txt` and the `scripts/` pipeline; per-structure evaluation flags are released to support paired re-analysis. The corrected simulator, the audit scripts of §5.7 and the per-structure records behind Table 2a are included.

Ethics Declaration

This study uses no human subjects, animal subjects, or sensitive data. The MP-20 dataset is composed entirely of publicly available crystal structures from the Materials Project.

References

- Batatia, I., Kovács, D. P., Simm, G. N. C., Ortner, C. & Csányi, G. (2022). MACE: higher order equivariant message passing neural networks for fast and accurate force fields. *Advances in Neural Information Processing Systems*, Vol. 35.
- Boulton, A. & Louër, D. (2004). *J. Appl. Cryst.* **37**, 724–731. DOI: 10.1107/S0021889804014876.
- Chitturi, S. R., Ratner, D., Walroth, R. C., Thampy, V., Reed, E. J., Dunne, M., Tassone, C. J. & Stone, K. H. (2021). *J. Appl. Cryst.* **54**, 1799–1810. DOI: 10.1107/S1600576721010840.
- Choudhary, K. (2025). *J. Phys. Chem. Lett.* **16**, 2110–2119. DOI: 10.1021/acs.jpcclett.4c03137. Reproduced from HF checkpoint `knc6/diffractgpt_mistral_chemical_formula` and code at `github.com/atomgptlab/atomgpt`.
- Cuocci, C., Corriero, N., Dell’Aera, M., Falcicchio, A., Rizzi, R. & Altomare, A. (2022). *Comput. Mater. Sci.* **210**, 111465. DOI: 10.1016/j.commatsci.2022.111465.
- Dauter, Z. & Jaskolski, M. (2010). *J. Appl. Cryst.* **43**, 1150–1171. DOI: 10.1107/S0021889810026956.
- de la Flor, G., Kroumova, E., Hanson, R. M. & Aroyo, M. I. (2023). *J. Appl. Cryst.* **56**, 1824–1840. DOI: 10.1107/S1600576723009068.
- de Wolff, P. M. (1968). *J. Appl. Cryst.* **1**, 108–113. DOI: 10.1107/S002188986800508X.
- Favre-Nicolin, V. & Černý, R. (2002). *J. Appl. Cryst.* **35**, 734–743. DOI: 10.1107/S0021889802015236.
- Guo, G., Saidi, T. L., Terban, M. W., Valsecchi, M., Billinge, S. J. L. & Lipson, H. (2025). *Nat. Mater.* **24**, 1726–1734. DOI: 10.1038/s41563-025-02220-y (Author Correction: 10.1038/s41563-025-02301-y; preprint arXiv:2406.10796). Reproduced from HF checkpoint `therealgabeguo/cdvae_xrd_sinc100` and code at `github.com/gabeguo/cdvae_xrd`.
- Guo, J. & Schwalbe-Koda, D. (2026). Generative inversion of spectroscopic data for amorphous structure elucidation. arXiv:2603.23210.
- Jiao, R., Huang, W., Lin, P., Han, J., Chen, P., Lu, Y. & Liu, Y. (2023). Crystal structure prediction by joint equivariant diffusion. *Advances in Neural Information Processing Systems*, Vol. 36.
- Johansen, F. L., Friis-Jensen, U., Dam, E. B., Jensen, K. M. Ø., Mercado, R. & Selvan, R. (2025). deCIFer: crystal structure prediction from powder diffraction data using autoregressive language models. *Transactions on Machine Learning Research* (arXiv:2502.02189).
- Lai, Q., Xu, F., Yao, L., Gao, Z., Liu, S., Wang, H., Lu, S., He, D., Wang, L., Zhang, L., Wang, C. & Ke, G. (2025). *Adv. Sci.* **12**, 2410722. DOI: 10.1002/advs.202410722.
- Li, Q., Jiao, R., Wu, L., Zhu, T., Huang, W., Jin, S., Liu, Y., Weng, H. & Chen, X. (2025). *Nat. Commun.* **16**, 7428. DOI: 10.1038/s41467-025-62708-8.

Nichol, A. & Dhariwal, P. (2021). Improved denoising diffusion probabilistic models. *Proceedings of the 38th International Conference on Machine Learning*, PMLR **139**, 8162–8171.

Parackal, A. S., Goodall, R. E. A., Faber, F. A. & Armiento, R. (2024). *Phys. Rev. Mater.* **8**, 103801. DOI: 10.1103/PhysRevMaterials.8.103801.

Riesel, E. A., Mackey, T., Nilforoshan, H., Xu, M., Badding, C. K., Altman, A. B., Leskovec, J. & Freedman, D. E. (2024). *J. Am. Chem. Soc.* **146**, 30340–30348. DOI: 10.1021/jacs.4c10244. Code: github.com/ML-PXRD/Crystalize. We could not reproduce: the checkpoint download link is marked “not yet active” in the upstream README (verified 2026-06-01).

Schütt, K. T., Kindermans, P.-J., Saucedo, H. E., Chmiela, S., Tkatchenko, A. & Müller, K.-R. (2017). SchNet: a continuous-filter convolutional neural network for modeling quantum interactions. *Advances in Neural Information Processing Systems*, Vol. 30.

Segal, N., Subramanian, A., Li, M., Miller, B. K. & Gómez-Bombarelli, R. (2026). *Digital Discovery* **5**, 1590–1599. DOI: 10.1039/D6DD00017G. (Preprint: arXiv:2512.04036.)

Song, J., Meng, C. & Ermon, S. (2021). Denoising diffusion implicit models. *International Conference on Learning Representations*.

Toby, B. H. & Von Dreele, R. B. (2013). *J. Appl. Cryst.* **46**, 544–549. Code: github.com/AdvancedPhotonSource/

Werner, P.-E., Eriksson, L. & Westdahl, M. (1985). *J. Appl. Cryst.* **18**, 367–370. DOI: 10.1107/S0021889885010512.

Xie, T., Fu, X., Ganea, O.-E., Barzilay, R. & Jaakkola, T. (2022). Crystal diffusion variational autoencoder for periodic material generation. *International Conference on Learning Representations*.

Yu, D., Zhu, Z., Leng, F. & Zhu, Y. (2026). *Nat. Commun.* **17**, 3274. DOI: 10.1038/s41467-026-70035-9.

Zeni, C., Pinsler, R., Zügner, D., Fowler, A., Horton, M., Fu, X., Wang, Z., Shysheya, A., Crabbé, J., Ueda, S., Sordillo, R., Sun, L., Smith, J., Nguyen, B., Schulz, H., Lewis, S., Huang, C.-W., Lu, Z., Zhou, Y., Yang, H., Hao, H., Li, J., Yang, C., Li, W., Tomioka, R. & Xie, T. (2025). *Nature* **639**, 624–632. DOI: 10.1038/s41586-025-08628-5.

Author note on reproduced numbers. The reproduced PXRDnet and DiffractGPT numbers in §5.3 are computed by the author on the released checkpoints through our shared **StructureMatcher** harness, not transcribed from the cited papers. The central claims depend only on these reproduced numbers under one common evaluation protocol, with confidence intervals stated throughout.

Appendix

A. Full ablation history

For completeness, Table A1 lists every training run discussed in the development of this paper, including those that did not make it into the main ablation table. Logs and checkpoints for all runs are in the released repository under **runs/**.

Run	Parameters	What changed	Result	Status in paper
v4	3.5 M	Global PXRD pooling, additive conditioning	Coord loss flat at 3.0	§3.2
v5	3.7 M	+ Multi-resolution cross-attention	Coord 3.0 $\rightarrow$ 1.0	§3.2
v6–v9	3.7 M	λ_{Debye} sweep $\{0, 0.1, 1, 10\}$, ε -prediction	All ~ 1 % match (within noise)	§5.2
v10	3.7 M	+ Lattice-input fix, $\lambda_{\text{Debye}} = 0$	Lat loss 1.0 $\rightarrow$ 0.05	§5.2
v11	3.7 M	+ Lattice-input fix, $\lambda_{\text{Debye}} = 1$	Match 0.9 %	Table 1
v12	10.1 M	Larger model ($d=384$, $L=5$), ε -prediction	Killed at 18 k; same plateau	§5.2
v13	3.7 M	x_0 -residual + lat-fix + Debye $\lambda=1$	2.51 % match (Phase 4 best, true-lat)	Table 1
v14	3.8 M	v13 + Wyckoff + distance loss	0.80 % match	Table 1
v15	3.8 M	v13 + Wyckoff only	2.10 % match	Table 1
v16	3.7 M	v13 + distance loss only	1.80 % match	Table 1
v17	3.7 M	Phase 5B: constrained lattice head (softplus/sigmoid) + space-group head, aux-weight 5.0, +30 k warm-start steps	Lattice MAE unchanged ($a \approx 1.11 \text{ \AA}$); SG head 39 % top-1	§5.6
v18 – v20	3.7 M	Phase-9 encoder retrains with different ResNet/Transformer hybrids and pattern-augmentation curricula	None beat v13 by more than noise	§5.1

Run	Parameters	What changed	Result	Status in paper
v21	3.7 M	Phase 9 final: v13 architecture + Phase 9 retrain + indexer drop-in support (submitted version; trained with the §5.7 (a) and (b) defects)	0 % (learned head), 1.6 % (indexer, crystal system given), 4.5 % (true-lat oracle), all 3-seed	§5.7
v22	3.7 M	v21 configuration retrained from scratch with the corrected form factor and the corrected x_0 -mode Debye loss (§5.7)	1.9 % (learned head), 1.6 % (indexer, crystal system given), 5.5 % (true-lat oracle), all 3-seed	Tables 2, 2a, 3

B. Hyperparameter sensitivity

We did not perform a full hyperparameter sweep. Pilot experiments on `d_model` $\in \{128, 256, 384\}$ and `n_layers` $\in \{2, 3, 5\}$ showed (256, 3) as a reasonable Pareto point. The Debye loss weight was swept $\{0, 0.1, 1, 10\}$ in early ε -prediction runs (v6–v9) without measurable effect on match rate; we re-fixed it at 1.0 for the x_0 -residual runs by analogy.

C. Differentiable simulator validation details

Pearson correlations between the DiffPXRD module and `pymatgen.XRDCalculator` on the 256-bin coarse grid, first 1000 MP-20 test structures (per-structure values in `paper/submissions/JAC-R1/analysis/simul` script `revalidate_simulator.py`). “Submitted” is the bare Gaussian sum of the submitted version; “corrected” is the `pymatgen` convention of §3.3. The ≤ 5 range is the training configuration; the ≤ 10 range is the default of `scripts/04_verify_debye.py`, behind the 0.962 reported in the submitted version.

Reflection range	Expression	n	Mean	Median	Min	Fraction > 0.9
$ h , k , l \leq 5$	submitted	1000	0.952	0.969	0.23	0.907
$ h , k , l \leq 5$	corrected	1000	0.988	0.999	0.25	0.980
$ h , k , l \leq 5$	submitted	50	0.939	0.971	0.49	0.90
$ h , k , l \leq 5$	corrected	50	0.974	0.998	0.52	0.92

Reflection range	Expression	n	Mean	Median	Min	Fraction > 0.9
$ h , k , l \leq 10$	submitted	1000	0.960	0.972	0.54	–
$ h , k , l \leq 10$	corrected	1000	0.999	1.000	0.58	–
$ h , k , l \leq 10$	submitted	50	0.963	0.972	0.90	–
$ h , k , l \leq 10$	corrected	50	0.999	1.000	0.98	–

The corrected expression scores at least as high as the submitted one on every one of the 1000 structures. Gradients through the structure factor with respect to fractional coordinates were verified non-zero and finite. The regression test `tests/test_debye_form_factor.py` pins $f(0) = Z$ and the agreement with `pymatgen` on NaCl.

D. Reproducibility checklist

- ☒ Hyperparameters specified (§4.3)
- ☒ Datasets and splits specified (§4.1; canonical CDVAE MP-20)
- ☒ Evaluation protocol specified (§4.2)
- ☒ Random seed: single seed (42) per run; pilot variance ≈ 0.5 % match-rate absolute
- ☒ Compute environment: PyTorch 2.x, CUDA 12.x, single RTX 5090; Python 3.12
- ☒ Code, checkpoints, logs, and full training scripts released